

Assignment 1: From Dirty Data to Predictive Models

Due Date: Monday, October 6, 2025 at 12:00 AM (midnight, Eastern Time)

Policy: <https://columbia-coms4995.github.io/aml-fall2025/#policies>



Submission

- A runnable Jupyter/Colab Notebook (.ipynb).
- A PDF report (5+ pages w/ clear structure).



Objective

The goal of this assignment is to guide you through the process of building an end-to-end flow to turn data to predictive models, using what you have learned so far. You will:

- Clean and transform a real dataset that contains missing values, noisy data, and categorical variables.
- Engineer meaningful features to improve model performance.
- Train and compare 2 different modeling approaches:
 - Generative: Naive Bayes
 - Discriminative: Linear Regression
- Evaluate your models using appropriate metrics and visualizations.
- Reflect on your workflow and disclose any AI tools used.

By the end of this assignment, you shall have hands-on experience with data preprocessing, feature engineering, model training, basic evaluation, and interpretation.



Datasets

Choose one of the following datasets. Both datasets contain real-world challenges such as missing values and categorical variables.

Titanic Survival Dataset

Task: Predict whether a passenger survived or not.

Dataset link: [Kaggle Titanic Dataset](#)

Notes: Requires a Kaggle account.

Heart Disease Prediction Dataset

Task: Predict whether a patient has heart disease based on medical attributes.

Dataset link: [UCI Heart Disease Dataset](#)

 Steps

To resolve this supervised learning task, your assignment will consist of the following steps:

1. Data Cleaning

- a. Handle missing values (imputation, dropping, or flagging).
- b. Address noisy or inconsistent values.
- c. Justify your choices (why impute vs. drop?).

2. Feature Engineering

- a. Apply appropriate transformations (normalization, standardization, log-scaling).
- b. Encode categorical variables (one-hot encoding or other methods).
- c. Optionally, construct new features (e.g., ratios, group statistics).

3. Model Training

- a. Train a Naive Bayes classifier (BernoulliNB or GaussianNB depending on your dataset), and make sure to apply Laplace smoothing (add-a).
 - i. You should experiment with at least 2 different values of alpha (e.g., 1.0 vs. 0.01) and compare the results.
- b. Train a Linear Regression model.
 - i. Although Linear Regression is typically used for **regression tasks**, in this assignment we will also apply it to a binary classification problem with a probability threshold as 0.5.
 - ii. Reminder: To prevent overfitting and to explore the effect of regularization, you are encouraged to also try:
 - 1. Ridge Regression (L2 regularization)
 - 2. LASSO Regression (L1 regularization)
- c. Ensure fair comparison (same train/test split).

4. Model Evaluation

- a. [Required] Accuracy + Confusion Matrix. [Encouraged] Precision, Recall, F1-score. [Bonus] ROC + AUC
- b. Include at least one visualization (confusion matrix heatmap, ROC curve).
- c. Explain what happened when you used smoothing vs. without smoothing.

5. Report Writing (5-7 pages)

- a. Introduction: problem definition + dataset description.
- b. Data Cleaning: your steps, reasoning, and before/after examples.
- c. Feature Engineering: transformations and encodings applied.
- d. Model Comparison: training setup, evaluation results.
- e. Discussion: interpretation of results and limitations.
- f. **AI Tool Usage Disclosure:**
 - i. List AI tools used (e.g., ChatGPT, Gemini, Claude).
 - ii. Describe what these tools contributed (e.g., generating starter code, visualization, debugging).
 - iii. **Clarify which parts were your own contribution.**

Pointers & Hints

To help you get started as the 1st assignment:

- Libraries: `scikit-learn`, `pandas`, `numpy`, `matplotlib/seaborn`.
- Functions:
 - `train_test_split` for splitting data.
 - `classification_report`, `confusion_matrix` for evaluation.
 - `cross_val_score` for validation (optional).
- Models:
 - `sklearn.naive_bayes.BernoulliNB`
 - `sklearn.naive_bayes.GaussianNB`
 - `sklearn.linear_model.LinearRegression`
- Visualization:
 - Use `seaborn.heatmap` for confusion matrix.
 - [Optional] Use `sklearn.metrics.roc_curve` and `auc` for ROC plots.

Grading Rubric

Category	Percentage	Description
Data Cleaning & Transformation	20	Clear handling of missing, noisy, and inconsistent data with justification.
Feature Engineering	20	Correct application of transformations and encodings; creativity in constructing new features.
Model Implementation	20	Properly trained Naive Bayes and Linear Regression; fair comparison setup.
Evaluation & Visualization	15	Correct use of metrics; meaningful visualizations included.
Discussion & Interpretation	15	Thoughtful analysis of results, strengths/limitations, and clear explanations.
AI Tool Usage Disclosure	10	Transparent description of AI tool usage and personal contributions.

Bonus (up to +10%)

Bonus points are only awarded for well-explained analysis, NOT just running extra code.

- Discuss how the choice of regularization (and the parameter `alpha`) affects model performance and feature importance.
- Observe which coefficients are shrunk to zero and discuss what this means for feature importance in LASSO.
- Reflect on AI vs. human-generated code and insights.